

Building simple Redis data caches

Dirk Eddelbuettel
dirk@eddelbuettel.com

Submission for R/Finance 2014

Redis, an Open Source "No-SQL" in-memory database with disk persistence, has become a key piece of infrastructure for many technology firms.¹ It offers the ability to write large amounts of data, even in bursts, with the ability to serve a large of subscribers. However, use cases from Finance are less well known. This proposed talk aims to present a simple data caching setup suitable for trading and investment environments.

Redis is agnostic to operating systems and programming languages, and in many cases requires only basic tcp/ip networking to interface its backend as the **rredis** package by Lewis has long demonstrated. **Redis** clients can be written in just about every compiled or scripting languages, including R.

Use from *R* is both tempting and convenient as the **rredis** packages takes care of the required serialization of *R* objects for both reading and writing. However, this limits the ability to read / write from other (common) languages which do not have access to *R*'s serialization. Moreover, while **rredis** is convenient for tasks with relatively limited data size or update frequencies, its performance can become a limiting factor.

We introduce the (still evolving) **rhiredis** package which is built on top of the standard **Redis C** library **hiredis** using **Rcpp**. We discuss several key design decision, and demonstrate a use case with performance gains of two orders of magnitude enabling use in multi-user multi-display near-real-time data provision suitable for trading environments.

References

- Eddelbuettel D (2014). *rhiredis: R interface to Redis using the hiredis library*. URL <http://github.com/eddelbuettel/rhiredis>.
- Eddelbuettel D, François R (2014). *Rcpp: Seamless R and C++ Integration*. R package version 0.11.0, URL <http://CRAN.R-Project.org/package=Rcpp>.
- Lewis BW (2013). *rredis: Redis client for R*. R package version 1.6.8, URL <http://CRAN.R-Project.org/package=rredis>.
- Noordhuis P, Sanfilippo S, contributors (2014). *hiredis: Minimalistic C client for Redis >= 1.2*. URL <http://github.com/redis/hiredis>.

¹Wikipedia opens its article with "Redis is an open-source, networked, in-memory, key-value data store with optional durability. It is written in ANSI C. The development of Redis has been sponsored by Pivotal since May 2013; [1] before that, it was sponsored by VMware. [2][3] According to the monthly ranking by DB-Engines.com, Redis is the most popular key-value store. [4]"